

Introduction To R Programming

R objects and functions

by Martin Frigaard

Written: September 30 2021

Updated: November 24 2021

[Created using the "λέξις" theme](#)

R Programming

R is a versatile language for data wrangling, visualization, and modeling

Resources

Link to slides

<https://mjfrigaard.github.io/csuc-data-journalism/slides.html>

Link to exercises

<https://mjfrigaard.github.io/csuc-data-journalism/lessons-exercises.html>

The background of the slide features the R Project logo, which consists of a large, light gray 'R' with a blue 'R' overlaid on top of it. The text 'Getting Started' is centered over the logo.

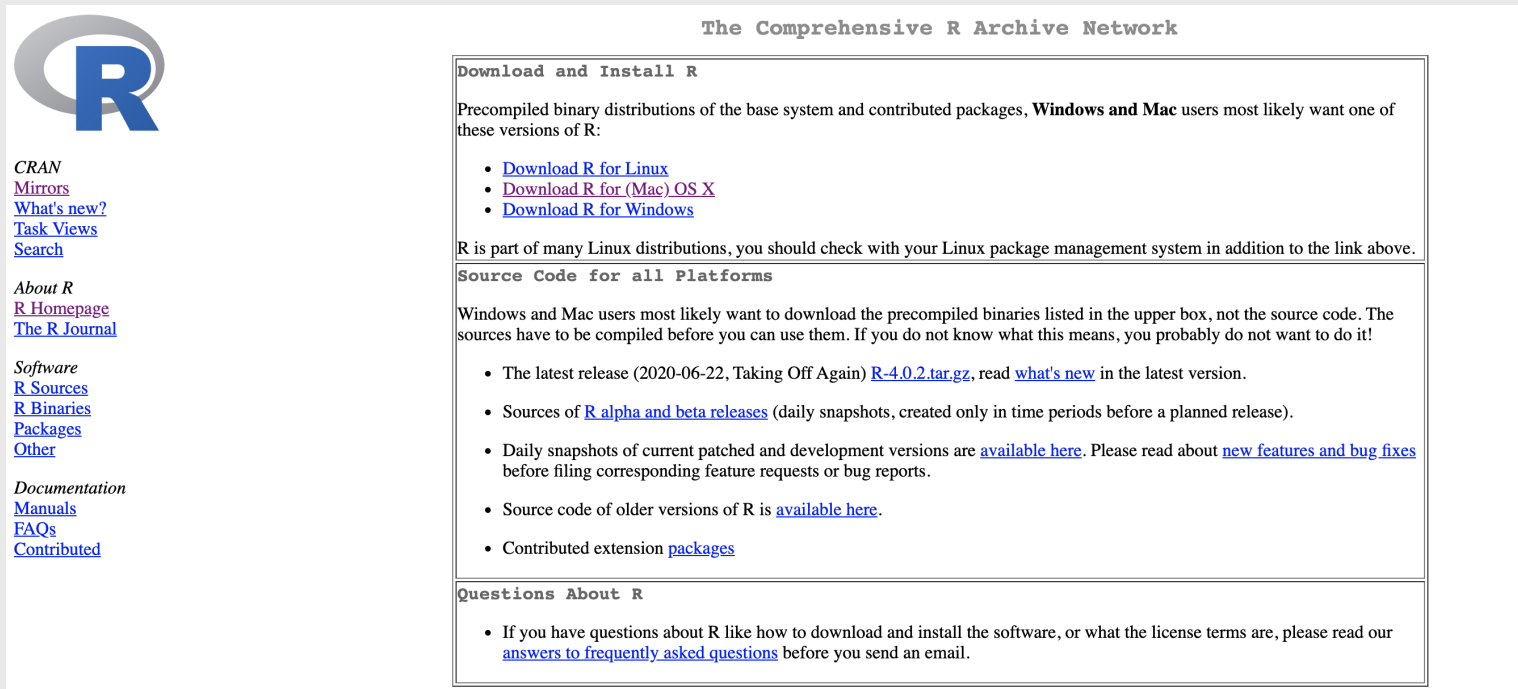
Getting Started

Image credit: [R Project](#)

Installing R

Install R from the Comprehensive R Archive Network (CRAN):

<https://cran.r-project.org/>



The screenshot displays the CRAN website. On the left is a navigation menu with links for CRAN, Mirrors, What's new?, Task Views, Search, About R, R Homepage, The R Journal, Software, R Sources, R Binaries, Packages, Other, Documentation, Manuals, FAQs, and Contributed. The main content area is titled 'The Comprehensive R Archive Network' and contains three sections: 'Download and Install R', 'Source Code for all Platforms', and 'Questions About R'. The 'Download and Install R' section provides precompiled binary distributions for Windows and Mac, with links to download R for Linux, Mac OS X, and Windows. The 'Source Code for all Platforms' section explains that Windows and Mac users should download precompiled binaries rather than source code, and lists links for the latest release, alpha and beta releases, daily snapshots, and source code of older versions. The 'Questions About R' section provides a link to frequently asked questions.

The Comprehensive R Archive Network

Download and Install R

Precompiled binary distributions of the base system and contributed packages, **Windows and Mac** users most likely want one of these versions of R:

- [Download R for Linux](#)
- [Download R for \(Mac\) OS X](#)
- [Download R for Windows](#)

R is part of many Linux distributions, you should check with your Linux package management system in addition to the link above.

Source Code for all Platforms

Windows and Mac users most likely want to download the precompiled binaries listed in the upper box, not the source code. The sources have to be compiled before you can use them. If you do not know what this means, you probably do not want to do it!










- The latest release (2020-06-22, Taking Off Again) [R-4.0.2.tar.gz](#), read [what's new](#) in the latest version.
- Sources of [R alpha and beta releases](#) (daily snapshots, created only in time periods before a planned release).
- Daily snapshots of current patched and development versions are [available here](#). Please read about [new features and bug fixes](#) before filing corresponding feature requests or bug reports.
- Source code of older versions of R is [available here](#).
- Contributed extension [packages](#)

Questions About R

- If you have questions about R like how to download and install the software, or what the license terms are, please read our [answers to frequently asked questions](#) before you send an email.


Download RStudio

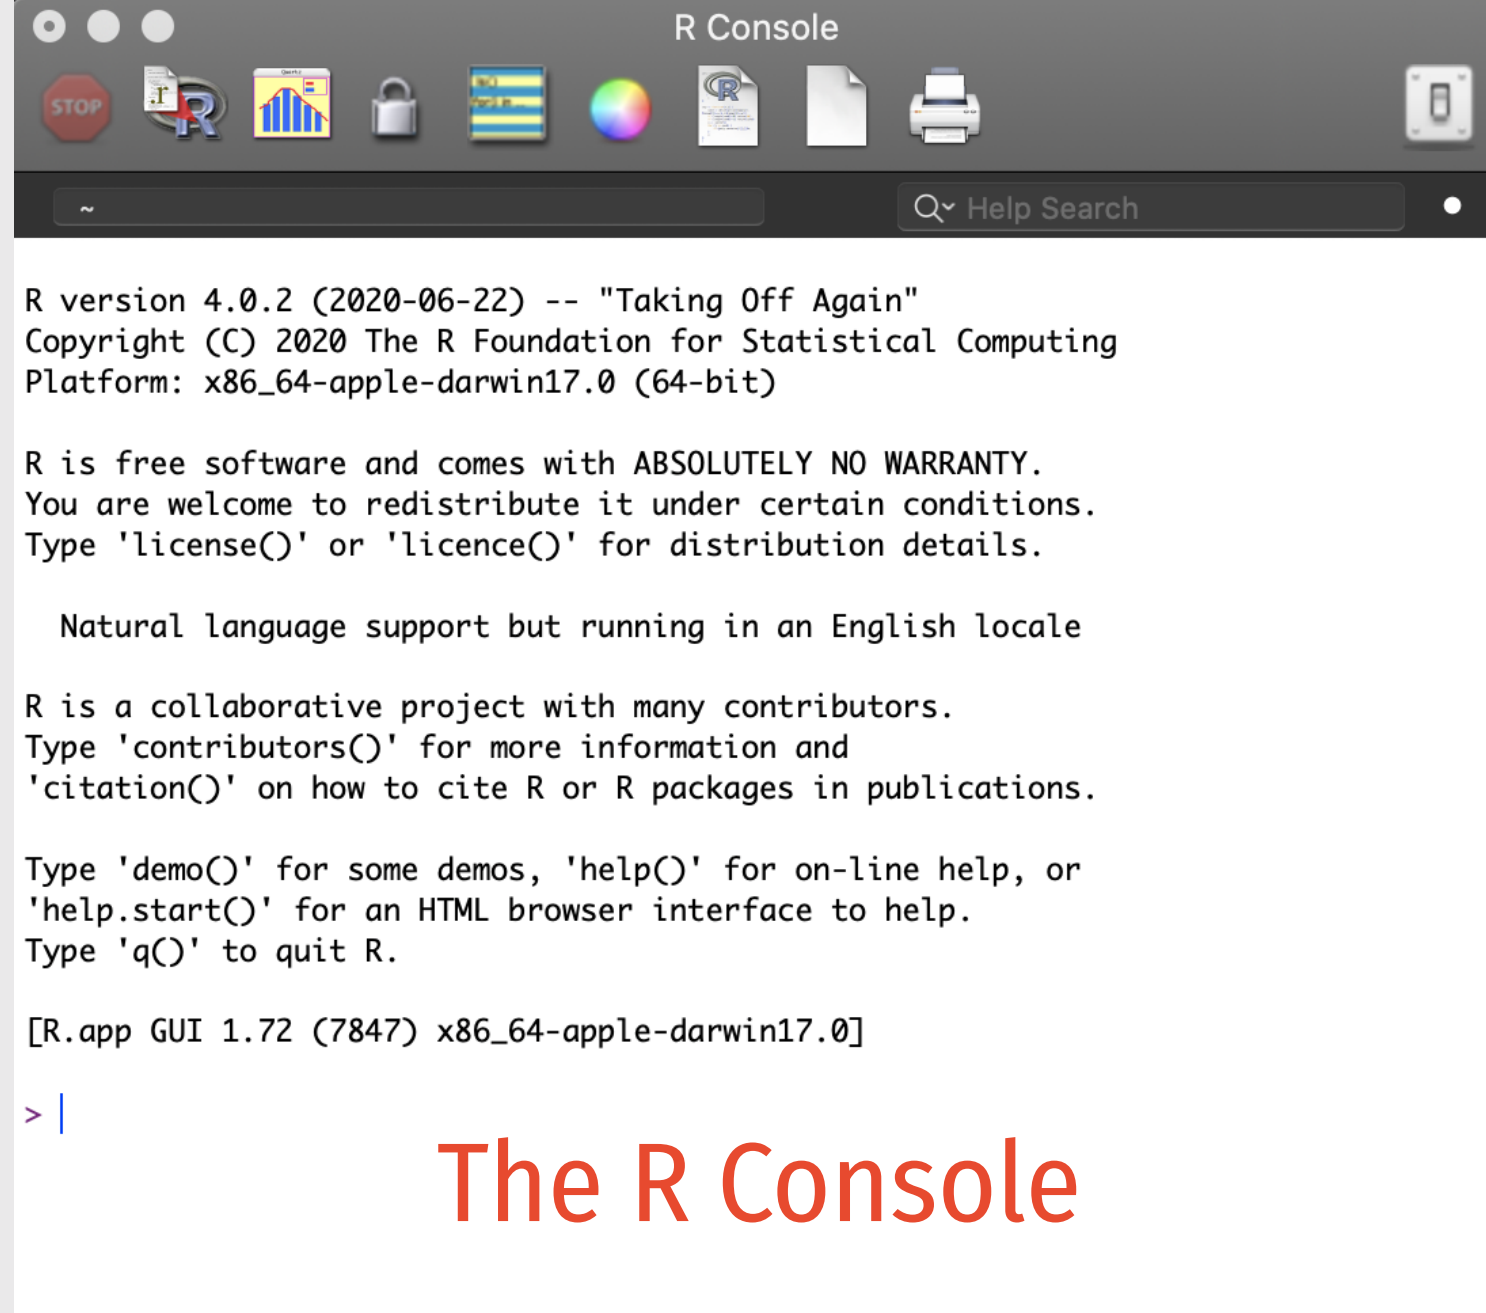
<https://rstudio.com/products/rstudio/download/>

OS	Download	Size	SHA-256
Windows 10/8/7	 RStudio-1.3.1093.exe	171.62 MB	62b9e60a
macOS 10.13+	 RStudio-1.3.1093.dmg	148.66 MB	bdc4d3a4
Ubuntu 16	 rstudio-1.3.1093-amd64.deb	124.33 MB	72f05048
Ubuntu 18/Debian 10	 rstudio-1.3.1093-amd64.deb	126.80 MB	ff222177
Fedora 19/Red Hat 7	 rstudio-1.3.1093-x86_64.rpm	146.96 MB	ed1f6ef8
Fedora 28/Red Hat 8	 rstudio-1.3.1093-x86_64.rpm	151.05 MB	01a978f3
Debian 9	 rstudio-1.3.1093-amd64.deb	127.00 MB	a747f9f9
SLES/OpenSUSE 12	 rstudio-1.3.1093-x86_64.rpm	119.43 MB	5016cbcf
OpenSUSE 15	 rstudio-1.3.1093-x86_64.rpm	128.40 MB	cf47e32d

Or use RStudio.Cloud

<https://rstudio.cloud/>





The R Console

Go to file/function

Addins

Untitled1

Source on Save

Run

Source

The RStudio IDE

1:1 (Top Level)

R Script

Environment

History

Connections

Tutorial

Import Dataset

List

R

Global Environment

Environment is empty

Console

Terminal

Jobs

~/R/

R version 4.0.2 (2020-06-22) -- "Taking Off Again"
Copyright (C) 2020 The R Foundation for Statistical Computing
Platform: x86_64-apple-darwin17.0 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

>

Files

Plots

Packages

Help

Viewer

Home

Find in Topic

Refresh Help Topic

R Resources

[Learning R Online](#)
[CRAN Task Views](#)
[R on StackOverflow](#)
[Getting Help with R](#)

RStudio

[RStudio IDE Support](#)
[RStudio Community Forum](#)
[RStudio Cheat Sheets](#)
[RStudio Tip of the Day](#)
[RStudio Packages](#)
[RStudio Products](#)

Manuals

9 / 60

<https://github.com/mjfrigaard/csuc-data-journalism>

Running R Commands

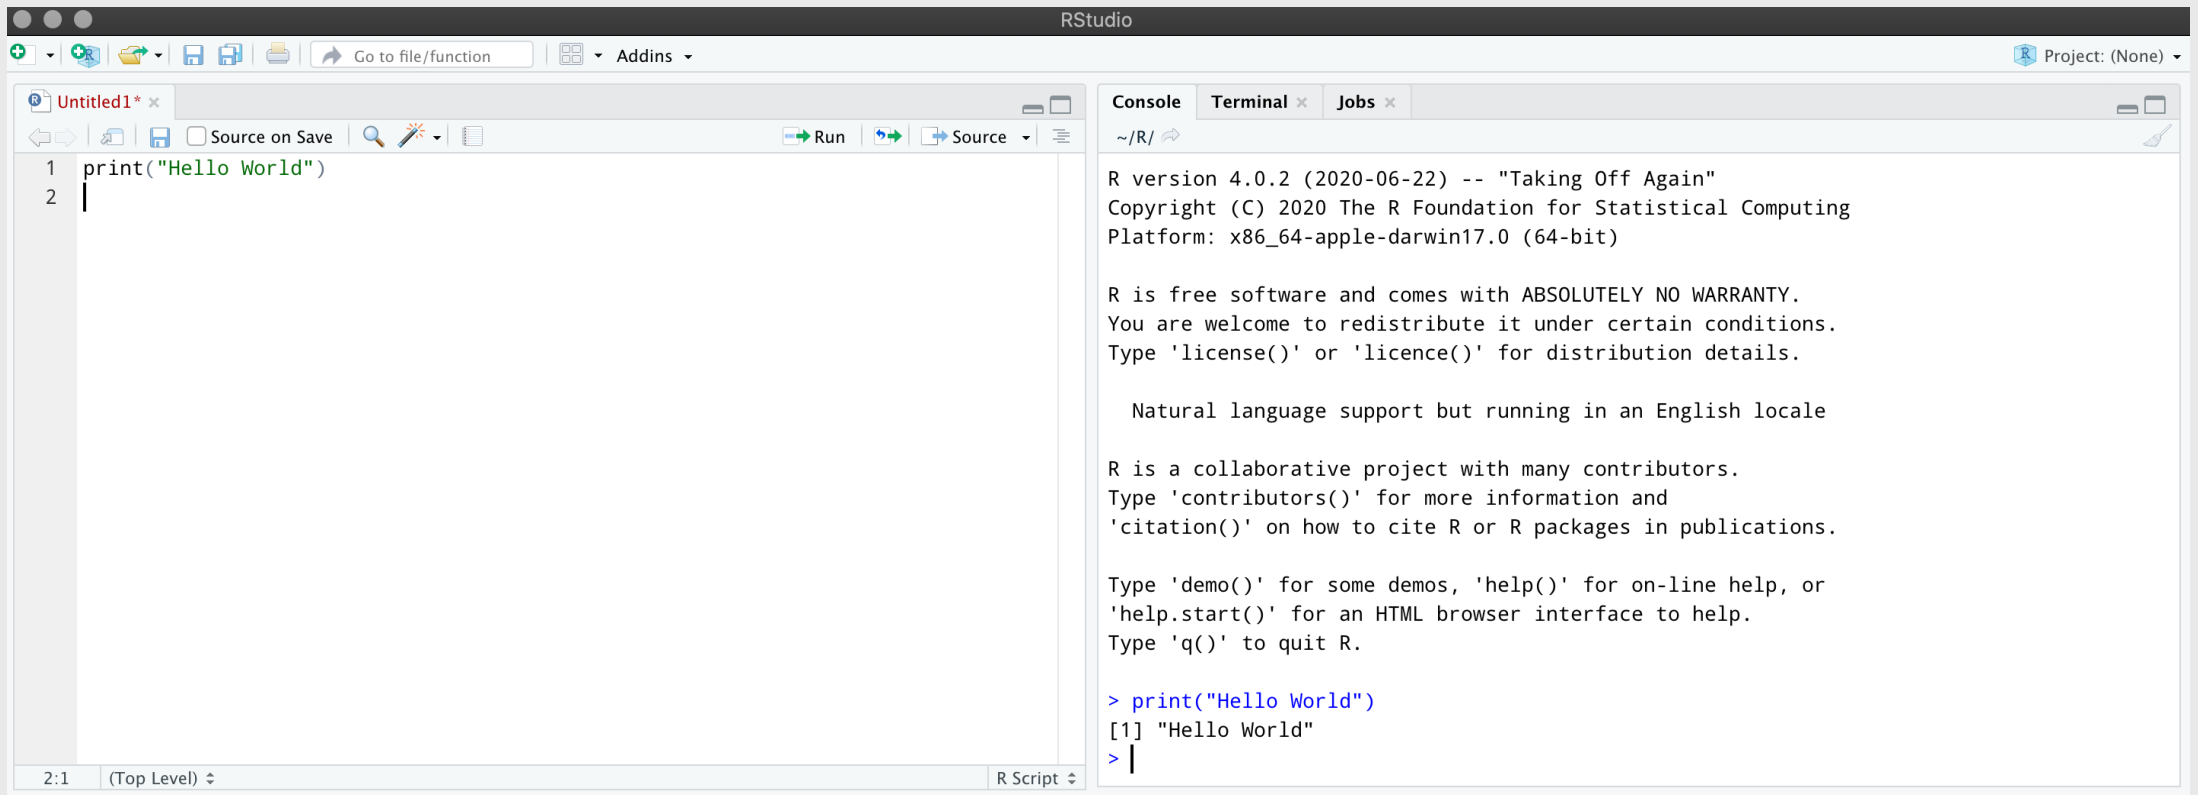
You can run R commands in the Console by entering them after the `>` operator (see example in R below)

```
print("Hello World")
```

```
[1] "Hello World"
```

Running R Commands

You can also run them in R scripts (see example in RStudio below)



The screenshot shows the RStudio application window. The top toolbar includes icons for file operations and a search bar. The main editor pane on the left shows a script named 'Untitled1*' with two lines of code: `1 print("Hello World")` and `2 |`. The bottom status bar indicates the cursor is at line 2, column 1, at the top level of an R script. The right-hand pane contains the 'Console' tab, which displays the R version information (4.0.2, 2020-06-22) and the output of the `print("Hello World")` command: `[1] "Hello World"`. The console also shows the prompt `> |` for the next command.

```
1 print("Hello World")
2 |
```

```
R version 4.0.2 (2020-06-22) -- "Taking Off Again"
Copyright (C) 2020 The R Foundation for Statistical Computing
Platform: x86_64-apple-darwin17.0 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> print("Hello World")
[1] "Hello World"
> |
```

R Syntax

The R syntax is comprised of two major elements:

Functions

Functions perform operations: calculate a mean, build a table, create a graph, etc.

Objects

Objects hold information: a collection of numbers, dates, words, models results, etc.

We use **functions** to perform operations on
objects

Example: create a vector of numbers

The standard assignment operator in R is `<-`. We can use this in combination with `c()` to create an object `x`, which contains five numbers (1, 3, 5, 7, 9).

```
x <- c(1, 3, 5, 7, 9)
```

Place `x` inside `print()` to print `x` to the console

```
x <- c(1, 3, 5, 7, 9)  
print(x)
```

NOTE: We can also use the `=` and move `->` to the end of the expression, but this is not recommended

R Syntax: functions

```
x <- c(1, 3, 5, 7, 9)
print(x)
```

```
[1] 1 3 5 7 9
```

In the example above, we've created object `x`, but what are `<-` and `c()`?

We can check this by passing them both in backticks to the `class()` function below.

```
class(`<-`)
```

```
[1] "function"
```

```
class(`c`)
```

```
[1] "function"
```

Functions in R

Functions perform operations (calculate, model, graph, etc.) on various **objects** that contain information (blood pressures, sales, political party affiliation, etc.)

Objects are similar to nouns: they hold information

```
object_1 <- "Sally"  
object_2 <- "dog"  
object_3 <- "road"
```

Functions are similar to verbs: they do things to nouns

```
work()  
run()  
implement()
```


Functions and objects

Functions perform operations on objects.

```
sally_object <- "Sally"  
work(sally_object)
```

Sally works.

```
dog_object <- "dog"  
run(dog_object)
```

The dog runs.

```
idea_object <- "idea"  
implement(idea_object)
```

Implement the idea.

Packages and functions in R

Functions are stored in R packages.

Fortunately, R comes 'out-of-the-box' with a set of functions for basic data management and statistical calculations.

To access the functions in a package, use the following syntax:

```
package::function(object)
```

The `median()` function comes from the `stats` package.

```
stats::median(x)
```

```
[1] 5
```

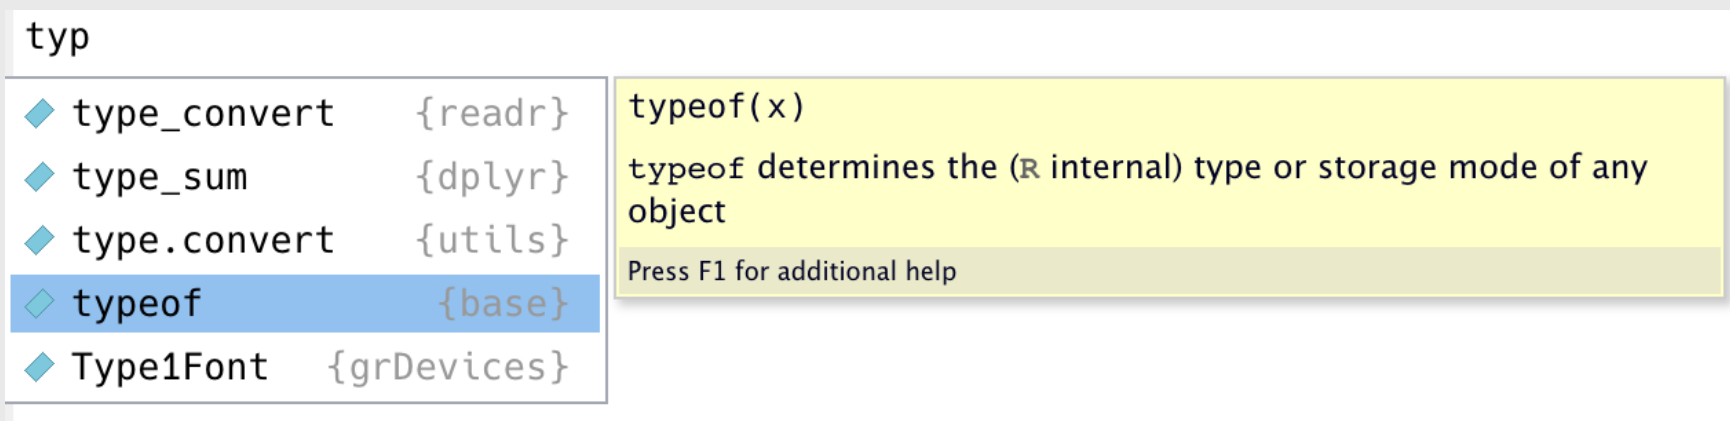
The `typeof()` function comes from the `base` package.

```
base::typeof(x)
```

```
[1] "double"
```

Packages and functions

Use tab-completion and the arrow keys in RStudio to explore a packages functions.



We can take advantage of tab-completion by using names that allow us to look up common objects. For example, naming plot objects with a `plot_` prefix will allow us to use tab-completion to scroll through each object without having to remember the specific name.

Installing packages from CRAN

To install packages from CRAN, we can use the `install.packages()` function.

```
install.packages("package name")
```


NOTE: if this is the first time installing packages, you'll probably be presented with a list of CRAN “mirrors” to use--choose the mirror closest to you.

To load the package into your environment, use `library(package name)`

```
library(package name)
```

Installing packages from CRAN in RStudio

You can also use the **Packages** pane in RStudio



The screenshot shows the RStudio interface with the 'Packages' pane active. The 'Install' button is highlighted with a red box. Below it, the 'Install Packages' dialog is open. The 'Install from:' dropdown is set to 'Repository (CRAN)'. The 'Packages (separate multiple with space or comma):' text box contains the placeholder text 'Package Name Here' and is also highlighted with a red box. The 'Install to Library:' dropdown is set to '/Library/Frameworks/R.framework/Versions/4.0/Resources/library'. The 'Install dependencies' checkbox is checked. At the bottom, there are 'Install' and 'Cancel' buttons.

Files	Plots	Packages	Help	Viewer
<input type="button" value="Install"/>	<input type="button" value="Update"/>			
Name	Description			

Install Packages

Install from: [? Configuring Repositories](#)

Repository (CRAN)

Packages (separate multiple with space or comma):

Package Name Here

Install to Library:

/Library/Frameworks/R.framework/Versions/4.0/Resources/library [C]

☒ Install dependencies

Installing user packages

The code for user-written packages are typically stored in code repository, like [Github](#).

To access user-written packages, you'll need to install the `devtools` or `remotes` packages.

```
install.packages("devtools")  
install.packages("remotes")
```

Use `devtools::install_github()` or `remotes::install_github()` (with the author's username and package repository name)

```
devtools::install_github(<username>/<package>)  
remotes::install_github(<username>/<package>)
```

Objects

R is typically referred to as an "object-oriented programming" language

We've covered functions, so now we'll dive into the aspects of some common R objects

Types of objects in R

- **Vectors**

- atomic (logical, integer, double, and character)
- S3 (factors, dates, date-times, durations)

- **Matrices**

- two dimensional objects

- **Arrays**

- multidimensional objects

- **Data frames & tibbles**

- rectangular objects

- **Lists**

- recursive objects

Atomic vectors

Vectors are the fundamental data type in R.

Many of R's functions are *vectorised*, which means they're designed for performing operations on vectors.

The "atomic" in atomic vectors means, "*of or forming a single irreducible unit or component in a larger system.*"

Atomic vectors can be logical, integer, double, or character (strings).


We will build each of these vectors using the previously covered assignment operator (`<—`) and `c()` function (*which stands for 'combine'*).

Store and explore

A common practice in R is to create an object, perform an operation on that object with a function, and store the results in new object.

We then explore the contents of the new object with another function.

Many of the functions in R are written with this *store and explore* process in mind.



Atomic vectors: numeric

The two atomic numeric vectors are integer and double.

Integer vectors are created with a number and capital letter **L** (i.e. **1L**, **10L**)

```
vec_integer <- c(1L, 10L, 100L)
```

Double vectors can be entered as decimals, but they can also be created in scientific notation (**2.46e8**), or values determined by the floating point standard (**Inf**, **-Inf** and **NaN**).

```
vec_double <- c(0.1, 1.0, 10.01)
```

Atomic vectors: numeric

We will use the `typeof()` and `is.numeric()` functions to explore the contents of `vec_integer` and `vec_double`.

```
typeof(vec_integer)
```

```
[1] "integer"
```

```
is.numeric(vec_integer)
```

```
[1] TRUE
```

`typeof()` tells us that this is an `"integer"` vector, and `is.numeric()` tests to see if it is numeric (which is `TRUE`).

Atomic vectors: logical vectors

Logical vectors can be `TRUE` or `FALSE` (or `T` or `F` for short). Below we use `typeof()` and `is.logical()` to explore the contents of `vec_logical`.

```
vec_logical <- c(TRUE, FALSE)
typeof(vec_logical)
```

```
[1] "logical"
```

```
is.logical(vec_logical)
```

```
[1] TRUE
```

Atomic vectors: logical vectors

Logical vectors are handy because when we add them together, and the total number tells us how many **TRUE** values there are.

```
TRUE + TRUE + FALSE + TRUE
```

```
[1] 3
```

Logical vectors can be useful for subsetting (a way of extracting certain elements from a particular object) based on a set of conditions.

*How many elements in **vec_integer** are greater than **5**?*

```
vec_integer > 5
```

```
[1] FALSE TRUE TRUE
```

Atomic vectors: character vectors

Character vectors store text data (note the double quotes). We'll *store and explore* again.

```
vec_character <- c("A", "B", "C")  
typeof(vec_character)
```

```
[1] "character"
```

```
is.character(vec_character)
```

```
[1] TRUE
```

Character vectors typically store text information that we need to include in a calculation, visualization, or model. In these cases, we'll need to convert them into **factors**. We'll cover those next.

S3 vectors

S3 vectors can be factors, dates, date-times, and difftimes.

S3 vectors: factors

Factors are categorical vectors with a given set of responses. Below we create a factor with three levels: `low`, `medium`, and `high`

```
vec_factor <- factor(x = c("low", "medium", "high"))  
class(vec_factor)
```

```
[1] "factor"
```

Factors are not character variables, though. They get stored with an integer indicator for each character level.

```
typeof(vec_factor)
```

```
[1] "integer"
```

S3 vectors: factor attributes

Factors are integer vectors with two additional attributes: `class` is set to `factor`, and `levels` for each unique response.

We can check this with `unique()` and `attributes()` functions.

```
unique(vec_factor)
```

```
[1] low    medium high  
Levels: high low medium
```

```
attributes(vec_factor)
```

```
$levels  
[1] "high"  "low"   "medium"  
  
$class  
[1] "factor"
```

S3 vectors: factor attributes

Levels are assigned alphabetically, but we can manually assign the order of factor levels with the `levels` argument in `factor()`.

```
vec_factor <- factor(  
  x = c("medium", "high", "low"),  
  levels = c("low", "medium", "high"))
```

We can check the levels with `levels()` or `unclass()`

```
levels(vec_factor)
```

```
[1] "low"      "medium" "high"
```

```
unclass(vec_factor)
```

```
[1] 2 3 1  
attr(,"levels")  
[1] "low"      "medium" "high"
```

S3 vectors: date

Dates are stored as `double` vectors with a `class` attribute set to `Date`.

R has a function for getting today's date, `Sys.Date()`. We'll create a `vec_date` using `Sys.Date()` and adding 1 and 2 to this value.

```
vec_date <- c(Sys.Date(),  
              Sys.Date() + 1,  
              Sys.Date() + 2)  
vec_date
```

```
[1] "2021-11-24" "2021-11-25" "2021-11-26"
```

We can see adding units to the `Sys.Date()` added days to today's date. The `attributes()` function tells us this vector has it's own class.

```
attributes(vec_date)
```

```
$class  
[1] "Date"
```

S3 vectors: date calculations

Dates are stored as a number because they represent the amount of days since January 1, 1970, which is referred to as the [UNIX Epoch](#).

`unclass()` tells us what the actual number is.

```
unclass(vec_date)
```

```
[1] 18955 18956 18957
```

S3 vectors: date-time

Date-times contain a bit more information than dates. The function to create a datetime vector is `as.POSIXct()`.

We'll convert `vec_date` to a date-time and store it in `vec_datetime_ct`. View the results below.

```
vec_date
```

```
[1] "2021-11-24" "2021-11-25" "2021-11-26"
```

```
vec_datetime_ct <- as.POSIXct(x = vec_date)
vec_datetime_ct
```

```
[1] "2021-11-23 17:00:00 MST" "2021-11-24 17:00:00 MST"
[3] "2021-11-25 17:00:00 MST"
```

We can see `vec_datetime_ct` stores some additional information.

S3 vectors: date-time attributes

`vec_datetime_ct` is a `double` vector with an additional attribute of `class` set to `"POSIXct" "POSIXt"`.

```
typeof(vec_datetime_ct)
```

```
[1] "double"
```


```
attributes(vec_datetime_ct)
```

```
$class  
[1] "POSIXct" "POSIXt"
```

S3 vectors: date-time help

Read more about date-times by entering the `as.POSIXct` function into the console preceded by a question mark.

```
?as.POSIXct
```



The screenshot shows an R console window with the following commands and output:

```
~/Dropbox/@working-projects/01-high-priority/@BioMarin/r-meetup-materials/
> vec_datetime_ct
[1] "2020-10-09 17:00:00 MST" "2020-10-10 17:00:00 MST"
[3] "2020-10-11 17:00:00 MST"
> typeof(vec_datetime_ct)
[1] "double"
> attributes(vec_datetime_ct)
$class
[1] "POSIXct" "POSIXt"

> unclass(vec_datetime_ct)
[1] 1602288000 1602374400 1602460800
> attributes(vec_datetime_ct)
$class
[1] "POSIXct" "POSIXt"

> ?as.POSIXct
> |
```

Below the console is the R help window for `as.POSIXct`. The window title is "R: Date-time Conversion Functions". The search bar contains `as.POSIXct`. The help text is as follows:

as.POSIX* (base) R Documentation

Date-time Conversion Functions

Description

Functions to manipulate objects of classes "POSIXlt" and "POSIXct" representing calendar dates and times.

Usage

```
as.POSIXct(x, tz = "", ...)
as.POSIXlt(x, tz = "", ...)
```

S3 method for class 'character'

```
as.POSIXlt(x, tz = "", format,
  tryFormats = c("%Y-%m-%d %H:%M:%OS",
    "%Y/%m/%d %H:%M:%OS",
    "%Y-%m-%d %H:%M",
    "%Y/%m/%d %H:%M",
    "%Y-%m-%d",
    "%Y/%m/%d"),
  optional = FALSE, ...)
```

Default S3 method:

```
as.POSIXlt(x, tz = "",
  optional = FALSE, ...)
```

S3 method for class 'numeric'

```
as.POSIXlt(x, tz = "", origin, ...)
```

S3 method for class 'POSIXlt'

```
as.double(x, ...)
```


S3 vectors: difftime

Difftimes are durations, so we to create them with `time_01` and `time_02`:

```
time_01 <- Sys.Date()
time_02 <- Sys.Date() + 10
time_01
```

```
[1] "2021-11-24"
```

```
time_02
```

```
[1] "2021-12-04"
```

Difftimes are stored as a `double` vector.

```
vec_difftime <- difftime(time_01,
                          time_02,
                          units = "days")
vec_difftime
```

```
Time difference of -10 days
```

```
typeof(vec_difftime)
```

```
[1] "double"
```

S3 vectors: difftime attributes

Difftimes are their own `class` and have a `units` attribute set to whatever we've specified in the `units` argument.

```
attributes(vec_difftime)
```

```
$class  
[1] "difftime"  
  
$units  
[1] "days"
```

We can see the actual number stored in the vector with `unclass()`

```
unclass(vec_difftime)
```

```
[1] -10  
attr(,"units")  
[1] "days"
```

Matrices

A matrix is several vectors stored together into two a two-dimensional object.

```
mat_data <- matrix(
  data = c(vec_double,
vec_integer),
  nrow = 3, ncol = 2,
  byrow = FALSE)
mat_data
```

```
      [,1] [,2]
[1,] 0.10   1
[2,] 1.00  10
[3,] 10.01 100
```

We can check the dimensions of `mat_data` with `dim()`.

```
dim(mat_data)
```

```
[1] 3 2
```

This is a three-column, two-row matrix.

Matrix positions

The output in the console tells us where each element is located in `mat_data`.

For example, if I want to get the `10` that's stored in `vec_integer`, I can use look at the output and use the indexes.

```
mat_data
```

```
      [,1] [,2]  
[1,] 0.10   1  
[2,] 1.00  10  
[3,] 10.01 100
```

By placing the index (`[2, 2]`) next to the object, I am telling R, "*only return the value in this position*".

```
mat_data[2, 2]
```

```
[1] 10
```

Arrays

Arrays are like matrices, but they can have more dimensions.


```
dat_array <- array(  
  data = c(1, 2, 3, 4, 5,  
           6, 7, 8, 9, 10,  
           11, 12, 13, 14,  
           15, 16, 17, 18),  
  dim = c(3, 3, 2))
```

dat_array

```
, , 1  
      [,1] [,2] [,3]  
[1,]     1     4     7  
[2,]     2     5     8  
[3,]     3     6     9  
  
, , 2  
      [,1] [,2] [,3]  
[1,]    10    13    16  
[2,]    11    14    17  
[3,]    12    15    18
```

Array layers

`dat_array` contains numbers 1 through 18 in three columns and three rows, stacked in two *layers*.



```
class(dat_array)
```

```
[1] "array"
```

```
class(mat_data)
```

```
[1] "matrix" "array"
```

Matrices are arrays, but arrays are not matrices

Data Frames

Data frames are rectangular data with rows and columns (or observations and variables).

```
DataFrame <- data.frame(  
  character = c("A", "B",  
"C"),  
  integer = c(0.1, 1.0,  
10.01),  
  logical = c(TRUE, FALSE,  
TRUE))
```

DataFrame

	character	integer	logical
1	A	0.10	TRUE
2	B	1.00	FALSE
3	C	10.01	TRUE

NOTE: `stringsAsFactors = FALSE` is not required as of R version 4.0.0.

Data Frames

Check the structure of the `data.frame` with `str()`

```
str(DataFrame)
```

```
'data.frame':    3 obs. of  3 variables:
 $ character: chr  "A" "B" "C"
 $ integer  : num  0.1 1 10
 $ logical   : logi  TRUE FALSE TRUE
```

`str()` gives us a transposed view of the `DataFrame` object, and tells us the dimensions of the object.

Tibbles

Tibbles are a special kind of `data.frame` (*they print better to the console and character vectors are never coerced into factors*).

```
Tibble <- tibble::tribble(  
  ~character, ~integer, ~logical,  
  "A",      0.1,      TRUE,  
  "B",       1,     FALSE,  
  "C",    10.01,      TRUE)
```

The syntax to build them is slightly different, too.

Tibble

```
# A tibble: 3 × 3  
  character integer logical  
  <chr>      <dbl> <lgl>  
1 A          0.1 TRUE  
2 B           1 FALSE  
3 C        10.0 TRUE
```

Tibbles

Check the structure of `Tibble`.


```
str(Tibble)
```

```
tibble [3 × 3] (S3:
tbl_df/tbl/data.frame)
 $ character: chr [1:3] "A" "B" "C"
 $ integer  : num [1:3] 0.1 1 10
 $ logical  : logi [1:3] TRUE FALSE TRUE
```

`str()` tells us `tibbles` are `S3` objects, with types `tbl_df`, `tbl`, and `data.frame`.

Data frames and tibbles

If you're importing spreadsheets, most of the work you'll do in R will be with rectangular data objects (i.e. `data.frames` and `tibbles`).



These are the common rectangular data storage object for tabular data in R

Data frames & tibbles

DataFrame

```
character integer logical
1      A      0.10     TRUE
2      B      1.00    FALSE
3      C     10.01     TRUE
```

the `data.frame` prints the column names and contents

Tibble

```
# A tibble: 3 × 3
character integer logical
<chr>      <dbl> <lgl>
1 A          0.1  TRUE
2 B           1  FALSE
3 C         10.0  TRUE
```

the `tibble` prints the column names, dimensions, formats, and contents

Data frames & tibbles

If we check the `type` of the `DataFrame` and `Tibble`...

```
typeof(DataFrame)
```

```
[1] "list"
```

```
typeof(Tibble)
```

```
[1] "list"
```

...we see they are `lists`

Data Frames & Tibbles

Both `data.frames` and `tibbles` are their own class,

```
class(DataFrame)
```

```
[1] "data.frame"
```

```
class(Tibble)
```

```
[1] "tbl_df"      "tbl"        "data.frame"
```

So we can think of `data.frames` and `tibbles` as special kinds of *rectangular* lists, made with different types of vectors, with each vector being of equal length.

Lists

Lists are special objects because they can contain all other objects (including other lists).

```
dat_list <- list(  
  "integer" = vec_integer,  
  "array" = dat_array,  
  "matrix data" = mat_data,  
  "data frame" = DataFrame,  
  "tibble" = Tibble)
```

Lists have a `names` attribute, which we've defined above in double quotes.

```
attributes(dat_list)
```

```
$names  
[1] "integer"      "array"  
"matrix data" "data frame"  "tibble"
```

List structure

If we check the structure of the `dat_list`, we see the structure of list, and the structure of the elements in the list.

```
str(dat_list)
```

```
List of 5
 $ integer      : int [1:3] 1 10 100
 $ array        : num [1:3, 1:3, 1:2] 1 2 3 4 5 6 7 8 9 10 ...
 $ matrix data: num [1:3, 1:2] 0.1 1 10 1 10 ...
 $ data frame :'data.frame':   3 obs. of  3 variables:
  ..$ character: chr [1:3] "A" "B" "C"
  ..$ integer   : num [1:3] 0.1 1 10
  ..$ logical   : logi [1:3] TRUE FALSE TRUE
 $ tibble       : tibble [3 × 3] (S3: tbl_df/tbl/data.frame)
  ..$ character: chr [1:3] "A" "B" "C"
  ..$ integer   : num [1:3] 0.1 1 10
  ..$ logical   : logi [1:3] TRUE FALSE TRUE
```


Recap

In R, two major elements: functions and objects.

- *functions are verbs, objects are nouns*

Packages: use `install.packages()` and `library()` to load functions from packages

- *or `devtools::install_github(<username>/<package>)` or `remotes::install_github(<username>/<package>)`*

The most common R object is a vector

- Atomic vectors: *logical, integer, double, or character (strings)*
- S3 vectors: *factors, dates, date-times, and difftimes*

Recap, cont.

More complicated data structures: matrices and arrays

- Matrix: *two-dimensional object*
- Array: *multidimensional object*

Rectangular data structures:

- *data.frames* & *tibbles* are special kinds of rectangular lists, which can hold different types of vectors, with each vector being of equal length

Catch-all data structures:

- *lists* can contain all other objects (including other lists)

More resources

Learn more about R objects in the help files or the following online texts:

1. [R for Data Science](#)
2. [Advanced R](#)
3. [Hands on Programming with R](#)
4. [R Language Definition](#)

THANK YOU!

Feedback

@mjfrigaard on Twitter and Github

mjfrigaard@gmail.com